

PyTorch Seminar

Day2. CIFAR10 dataset training (2)

Duhyeon Kim / May 2025

 PyTorch

PyTorch Seminar

Day2. CIFAR10 dataset training

Duhyeon Kim / May 2025

 PyTorch

Contents

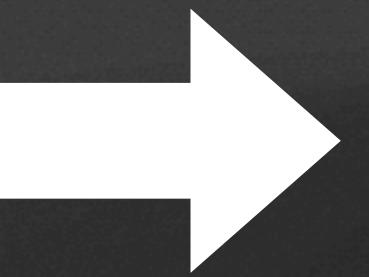
Things needed to Train CIFAR10

- A. Neural Network Composition
- B. CIFAR10 dataset
- C. Custom Dataloader (role of transform function)
- D. Custom Model (what is import ~ ?)

PyTorch nn composition

nn.Module

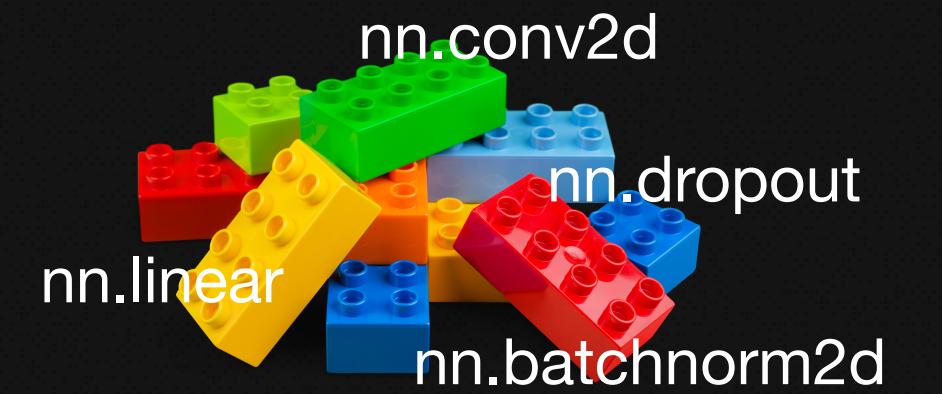
```
import torch.nn as nn
```



```
class CustomModel(nn.Module):
```

PyTorch nn composition

nn.Module

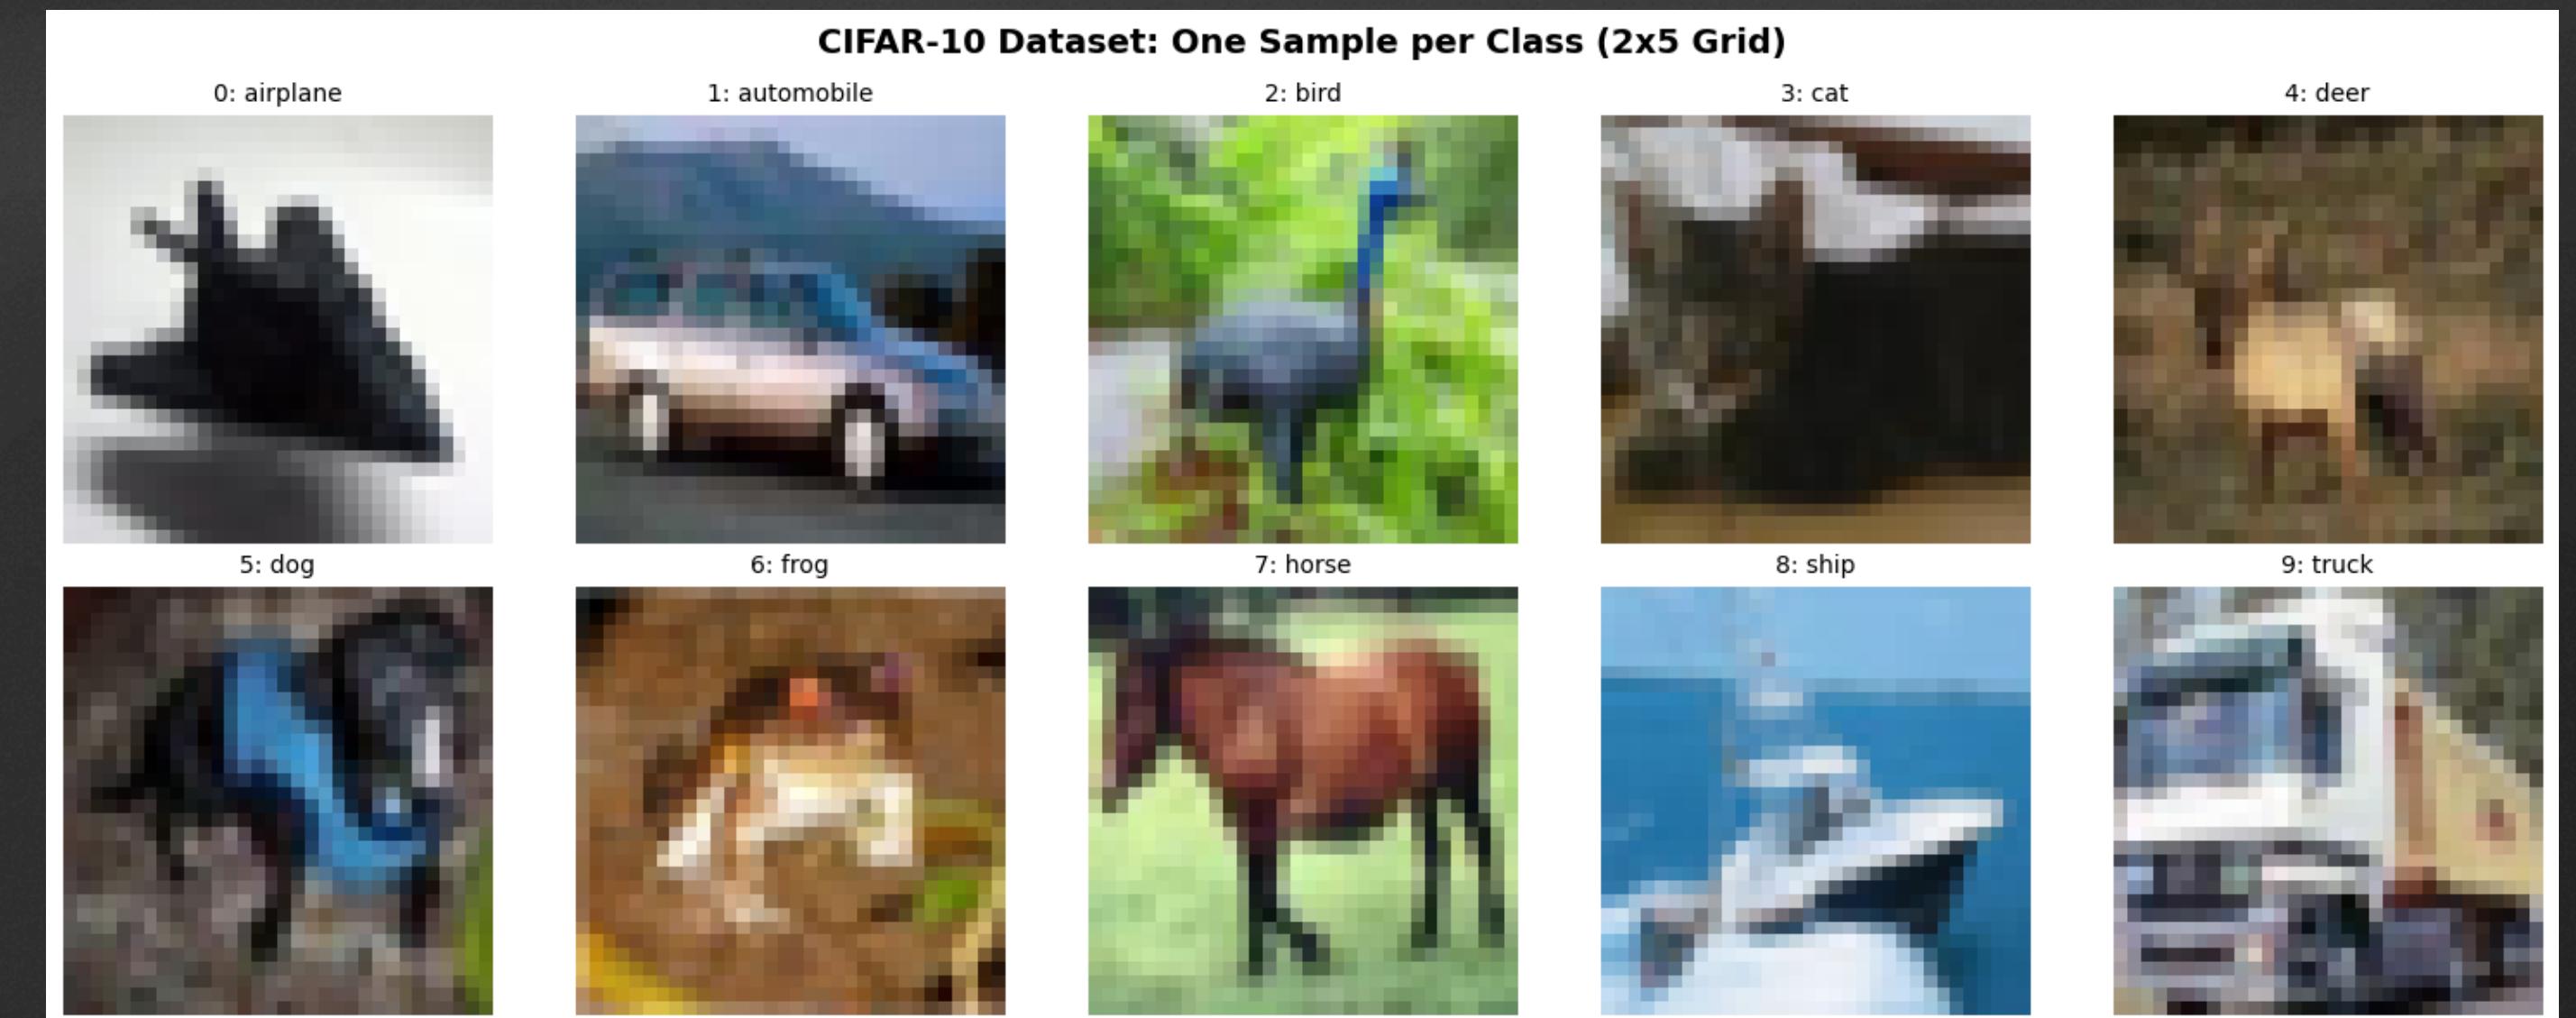


<https://colab.research.google.com/drive/16tkCZEnTU0bhtTvB0vBppk6TrO7IC47k?usp=sharing>

CIFAR-10 dataset

Overview

- 10 classes
{'airplane': 0, 'automobile': 1, 'bird': 2, 'cat': 3, 'deer': 4, 'dog': 5, 'frog': 6, 'horse': 7, 'ship': 8, 'truck': 9}
- 32x32 RGB images
- Train 50000 + Test 10000



<https://www.cs.toronto.edu/~kriz/cifar.html>

Simplest CIFAR-10 Training



Simplest CIFAR-10 Training

Evaluation Accuracy

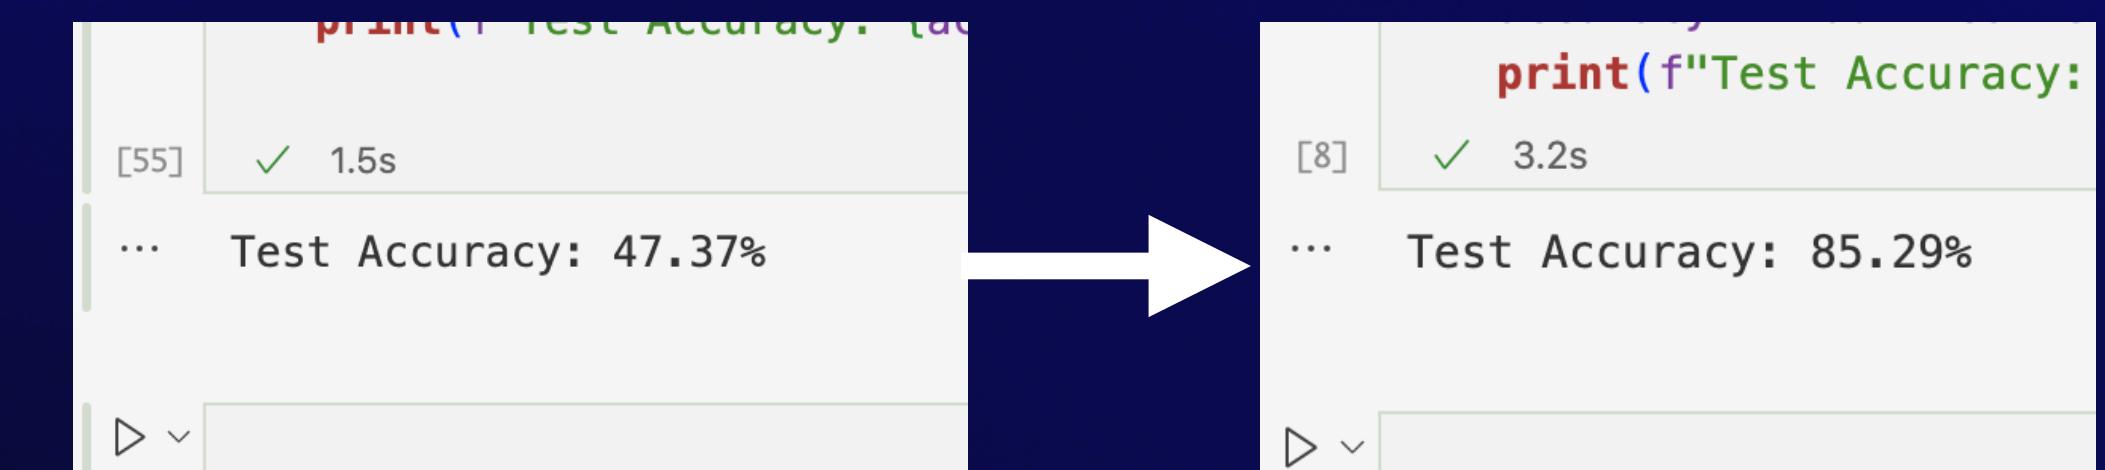
```
print('Test Accuracy: ', accuracy)

[55]:    ✓ 1.5s
... Test Accuracy: 47.37%
```

47% > 10% (random sampling)

How to further **improve** Model Performance?

Model Improvement Techniques



```
print(f"Test Accuracy: {accuracy:.2f}%")
```

[55] ✓ 1.5s
... Test Accuracy: 47.37%

▶ ▾

```
print(f"Test Accuracy: {accuracy:.2f}%")
```

[8] ✓ 3.2s
... Test Accuracy: 85.29%

▶ ▾

Model Improvement Techniques

Roads to 100% accur.

- A. Deeper Model
- B. Data Augmentation (torchvision.transform)
- C. Optimizer
- D. When it comes to LLM (Deepseek)

Deeper Model

AlexNet (Krizhevsky et al., 2012)

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca

Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca

Abstract

We trained a large, deep convolutional neural network to classify the 1.2 million high-resolution images in the ImageNet LSVRC-2010 contest into the 1000 different classes. On the test data, we achieved top-1 and top-5 error rates of 37.5% and 17.0% which is considerably better than the previous state-of-the-art. The neural network, which has 60 million parameters and 650,000 neurons, consists of five convolutional layers, some of which are followed by max-pooling layers, and three fully-connected layers with a final 1000-way softmax. To make training faster, we used non-saturating neurons and a very efficient GPU implementation of the convolution operation. To reduce overfitting in the fully-connected layers we employed a recently-developed regularization method called “dropout” that proved to be very effective. We also entered a variant of this model in the ILSVRC-2012 competition and achieved a winning top-5 test error rate of 15.3%, compared to 26.2% achieved by the second-best entry.

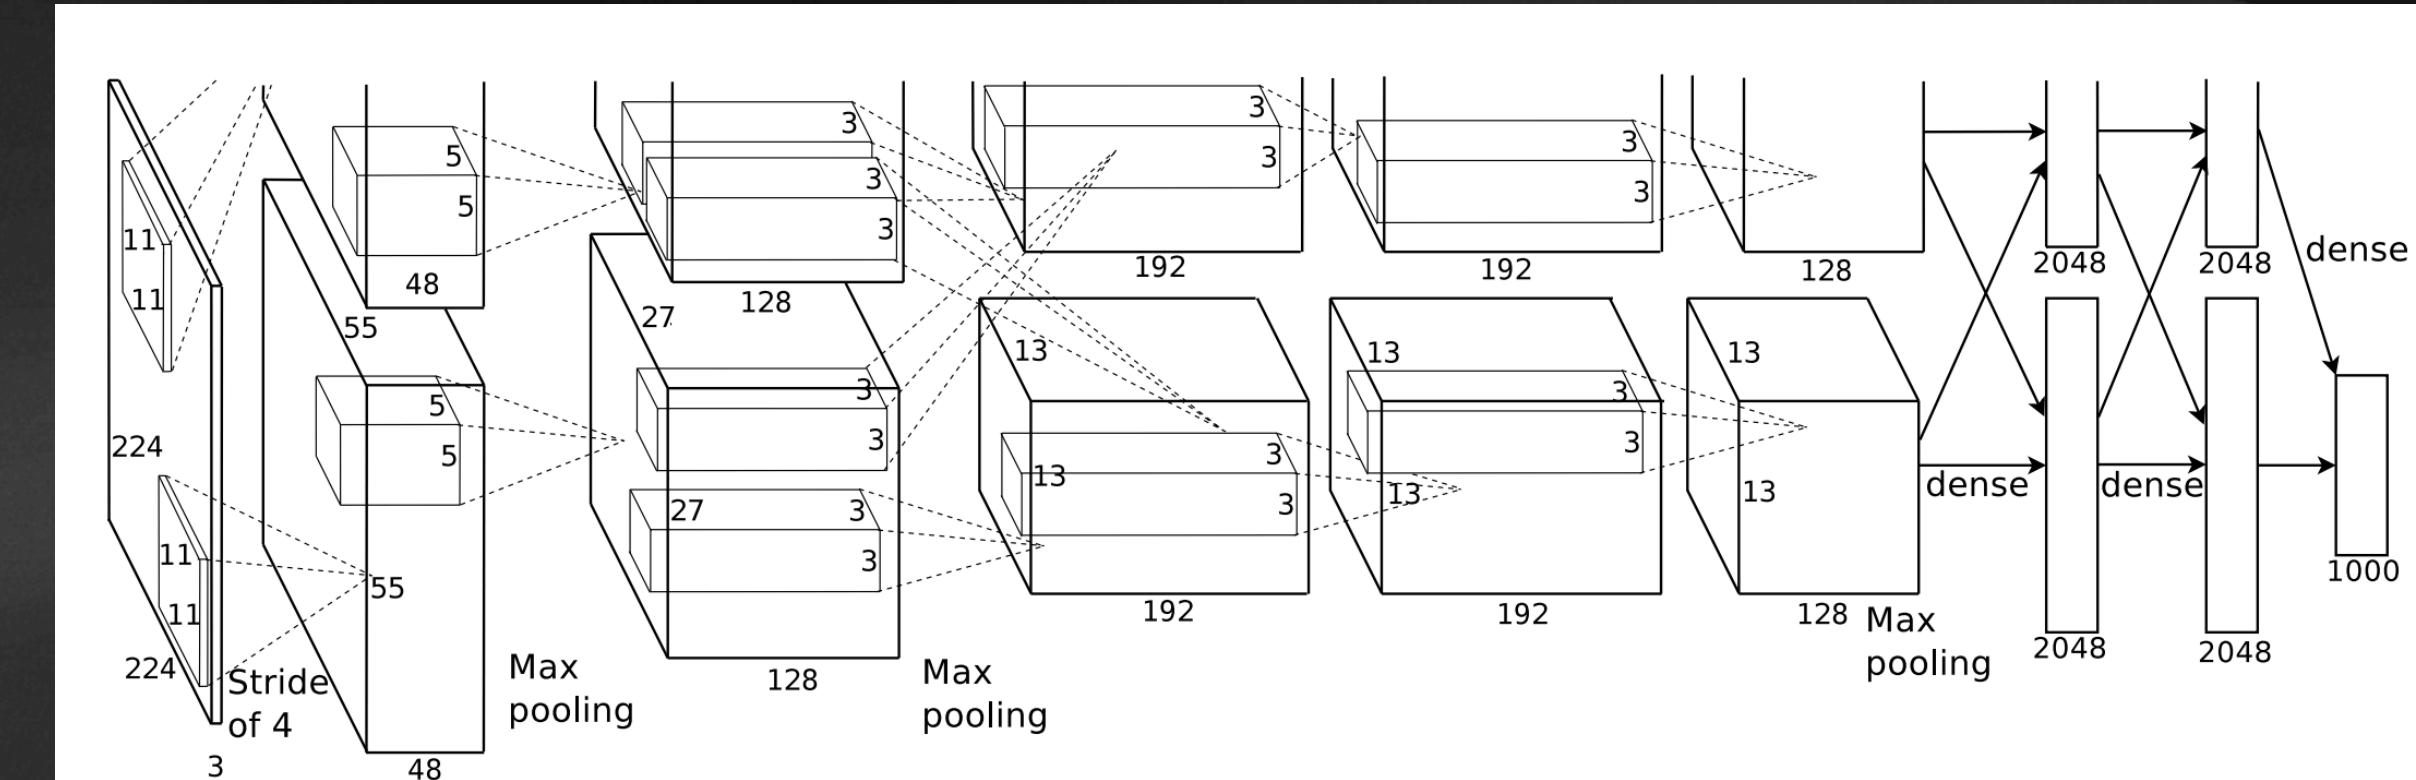


Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network’s input is 150,528-dimensional, and the number of neurons in the network’s remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Deeper Model

AlexNet (Krizhevsky et al., 2012)

Key Milestones

- **DEEP** Convolutional Network
 - 5 (conv) + 3 (FC)
- ReLU activation
 - Training Speed up (6x)
 - Gradient Vanishing solved
- GPU in practice
 - DEEP -> GPU

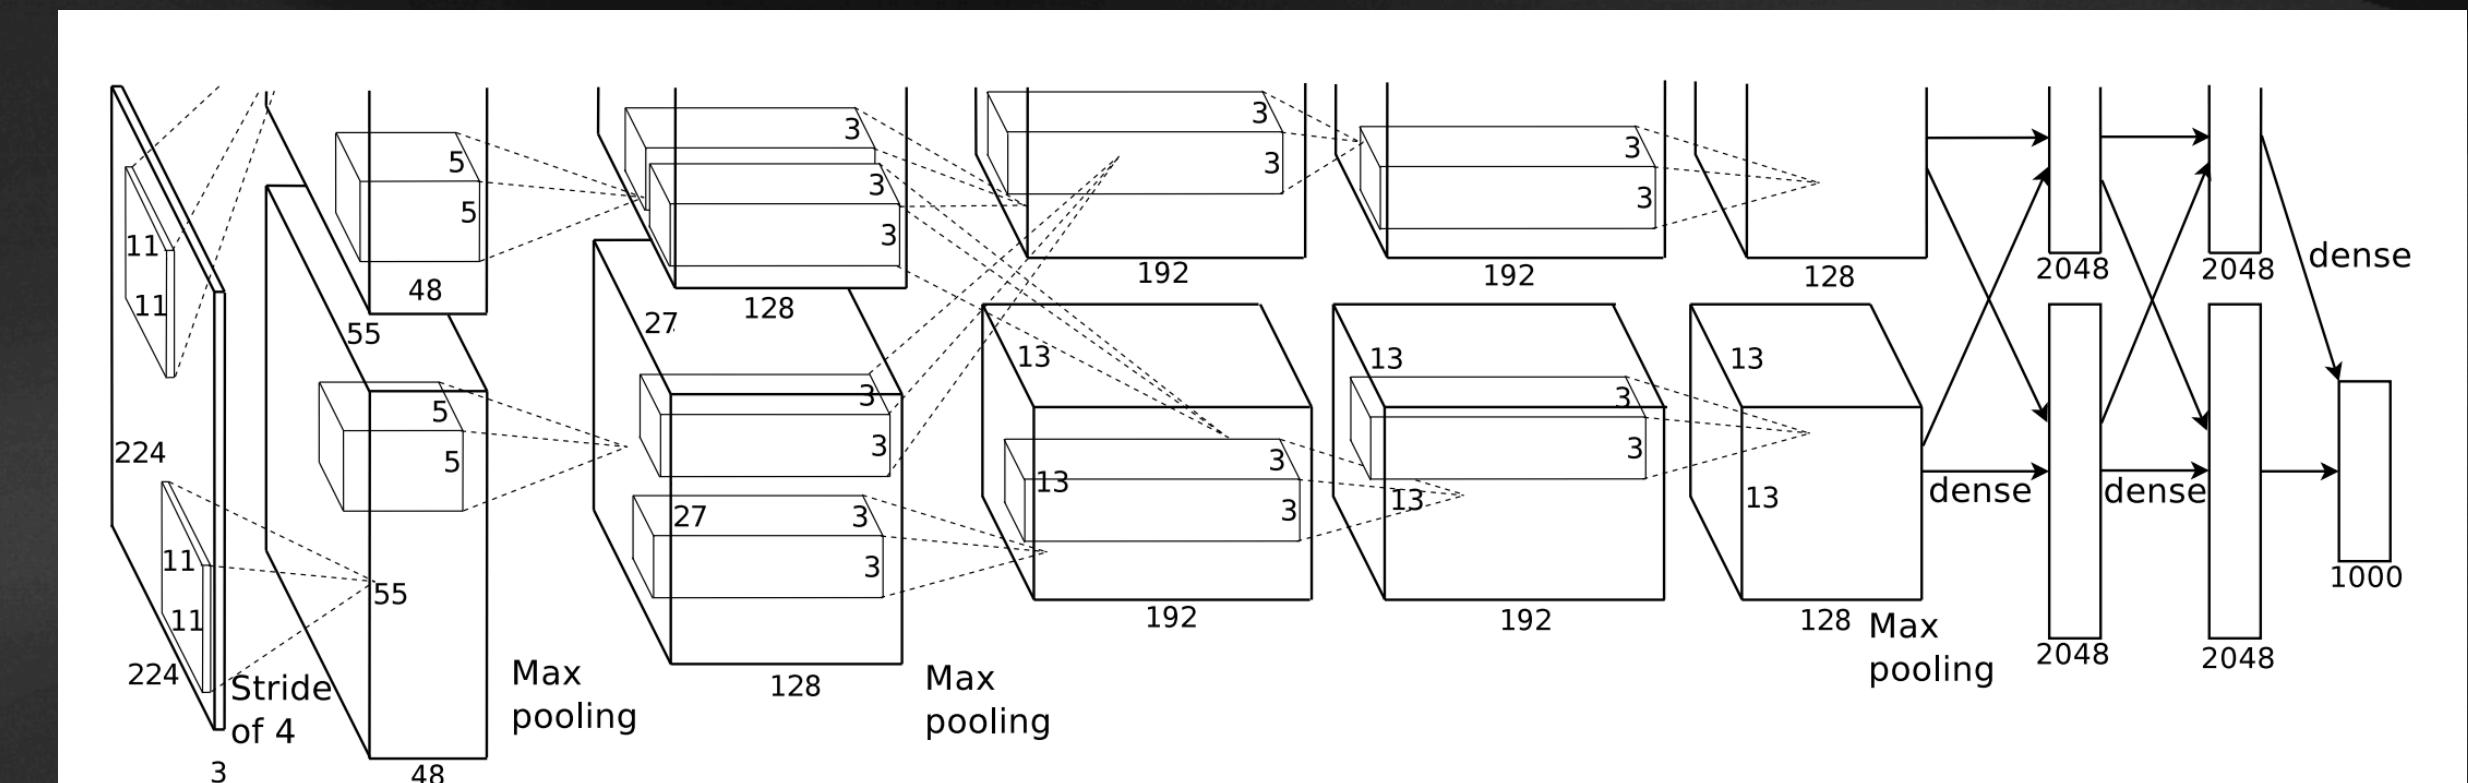
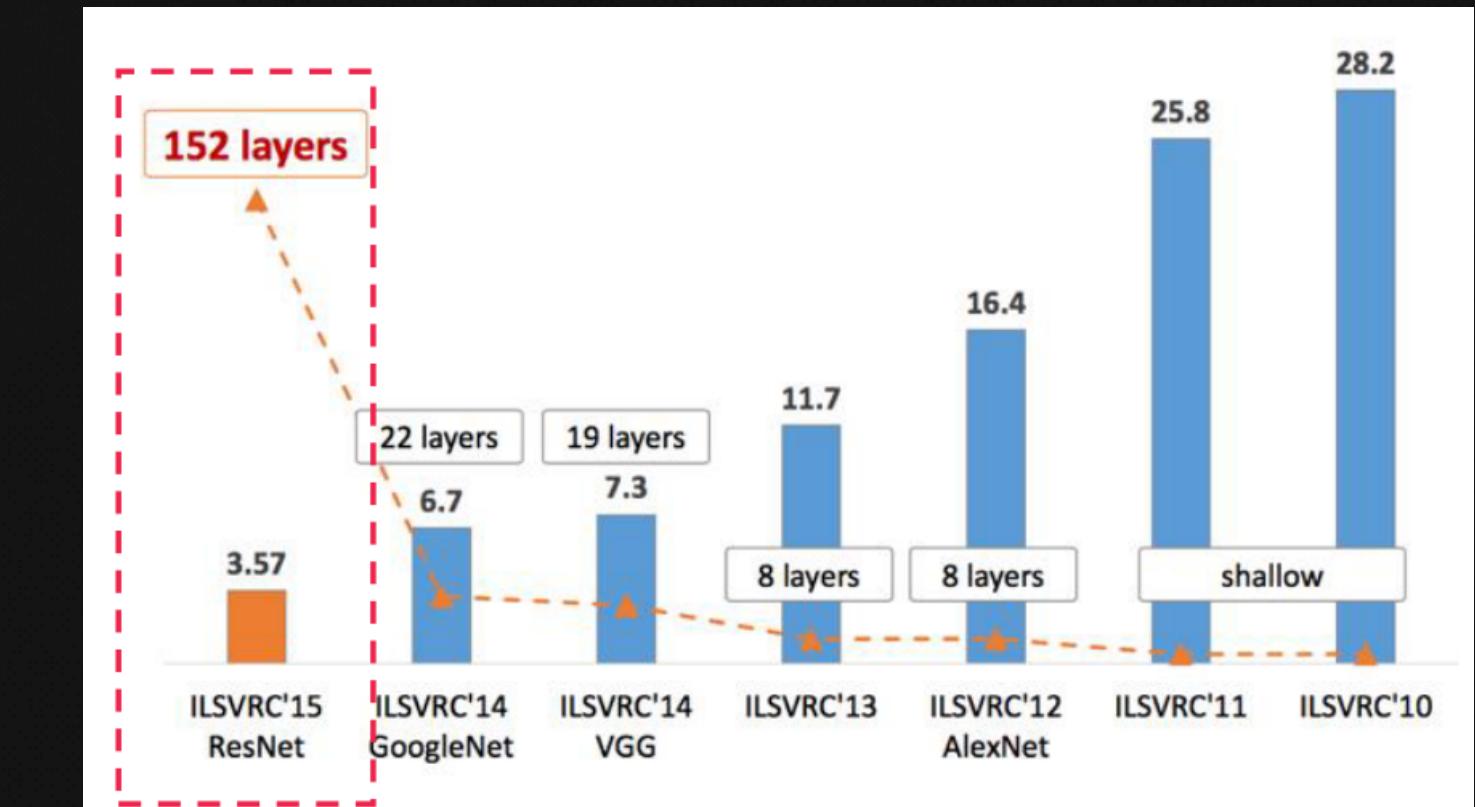


Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Data Augmentation

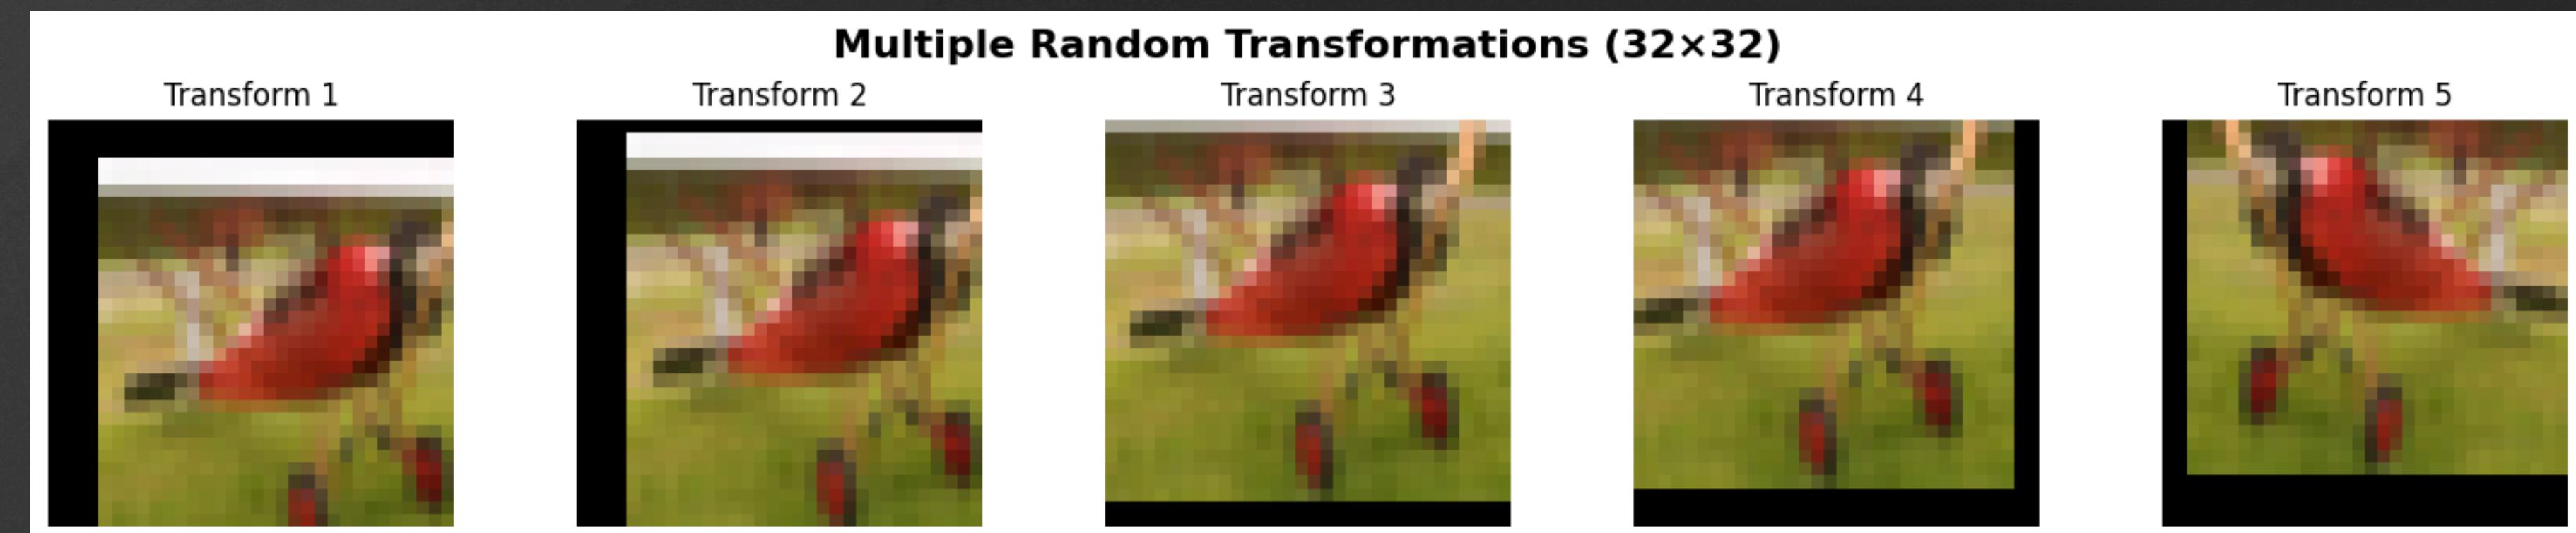
torchvision.transforms

Meaning
& Effect of
Augmenting Data



- increase dataset size and diversity
- Overfitting Prevention
 - sees diff view of the same image
- Data Efficiency
 - one datum -> several data

Performance Boost



Data Augmentation

torchvision.transforms

TORCHVISION.TRANSFORMS

Transforms are common image transformations. They can be chained together using `Compose`. Additionally, there is the `torchvision.transforms.functional` module. Functional transforms give fine-grained control over the transformations. This is useful if you have to build a more complex transformation pipeline (e.g. in the case of segmentation tasks).

All transformations accept PIL Image, Tensor Image or batch of Tensor Images as input. Tensor Image is a tensor with (C, H, W) shape, where C is a number of channels, H and W are image height and width. Batch of Tensor Images is a tensor of (B, C, H, W) shape, where B is a number of images in the batch. Deterministic or random transformations applied on the batch of Tensor Images identically transform all the images of the batch.

```
# MODIFIED: Enhanced data augmentation techniques and added normalization
train_transform = transforms.Compose([
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomCrop(32, padding=4), # MODIFIED: Added random crop
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
])
```

Transforms on `torch.*Tensor` only

Transform is used in dataset `__getitem__`

PIL image



transforms.Compose

RandomHorizontalFlip

RandomCrop

ToTensor

Normalize

Tensor image



Optimizer

SGD to Adam

Stochastic Gradient Descent

$$\theta_{t+1} = \theta_t - \eta \nabla L(\theta_t; x, y)$$

or

$$v_{t+1} = \beta v_t + \eta \nabla L(\theta_t; x, y)$$

$$\theta_{t+1} = \theta_t - v_{t+1}$$

Adaptive Moment Estimation

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) \nabla L(\theta_{t-1}) \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) (\nabla L(\theta_{t-1}))^2 \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t} \\ \hat{v}_t &= \frac{v_t}{1 - \beta_2^t} \end{aligned}$$

Momentum
Velocity

$$\theta_t = \theta_{t-1} - \frac{\eta \hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

Carefully and meaningfully update Model Parameter

Custom Model

Computer Vision

- Residual Network <https://arxiv.org/abs/1512.03385>
- 3x3 kernel <https://arxiv.org/abs/1409.1556>
- Generator and Discriminator <https://arxiv.org/abs/1406.2661>
- VAE <https://arxiv.org/abs/1312.6114>
- Feature Fusion https://openaccess.thecvf.com/content/WACV2021/papers/Dai_Attentional_Feature_Fusion_WACV_2021_paper.pdf

Custom Model

LLM (transformers)

- Multi-Head Attention
- Multi-Head latent Attention
- Multi-Query Attention
- Group-Query Attention

Memory + Inference Speed

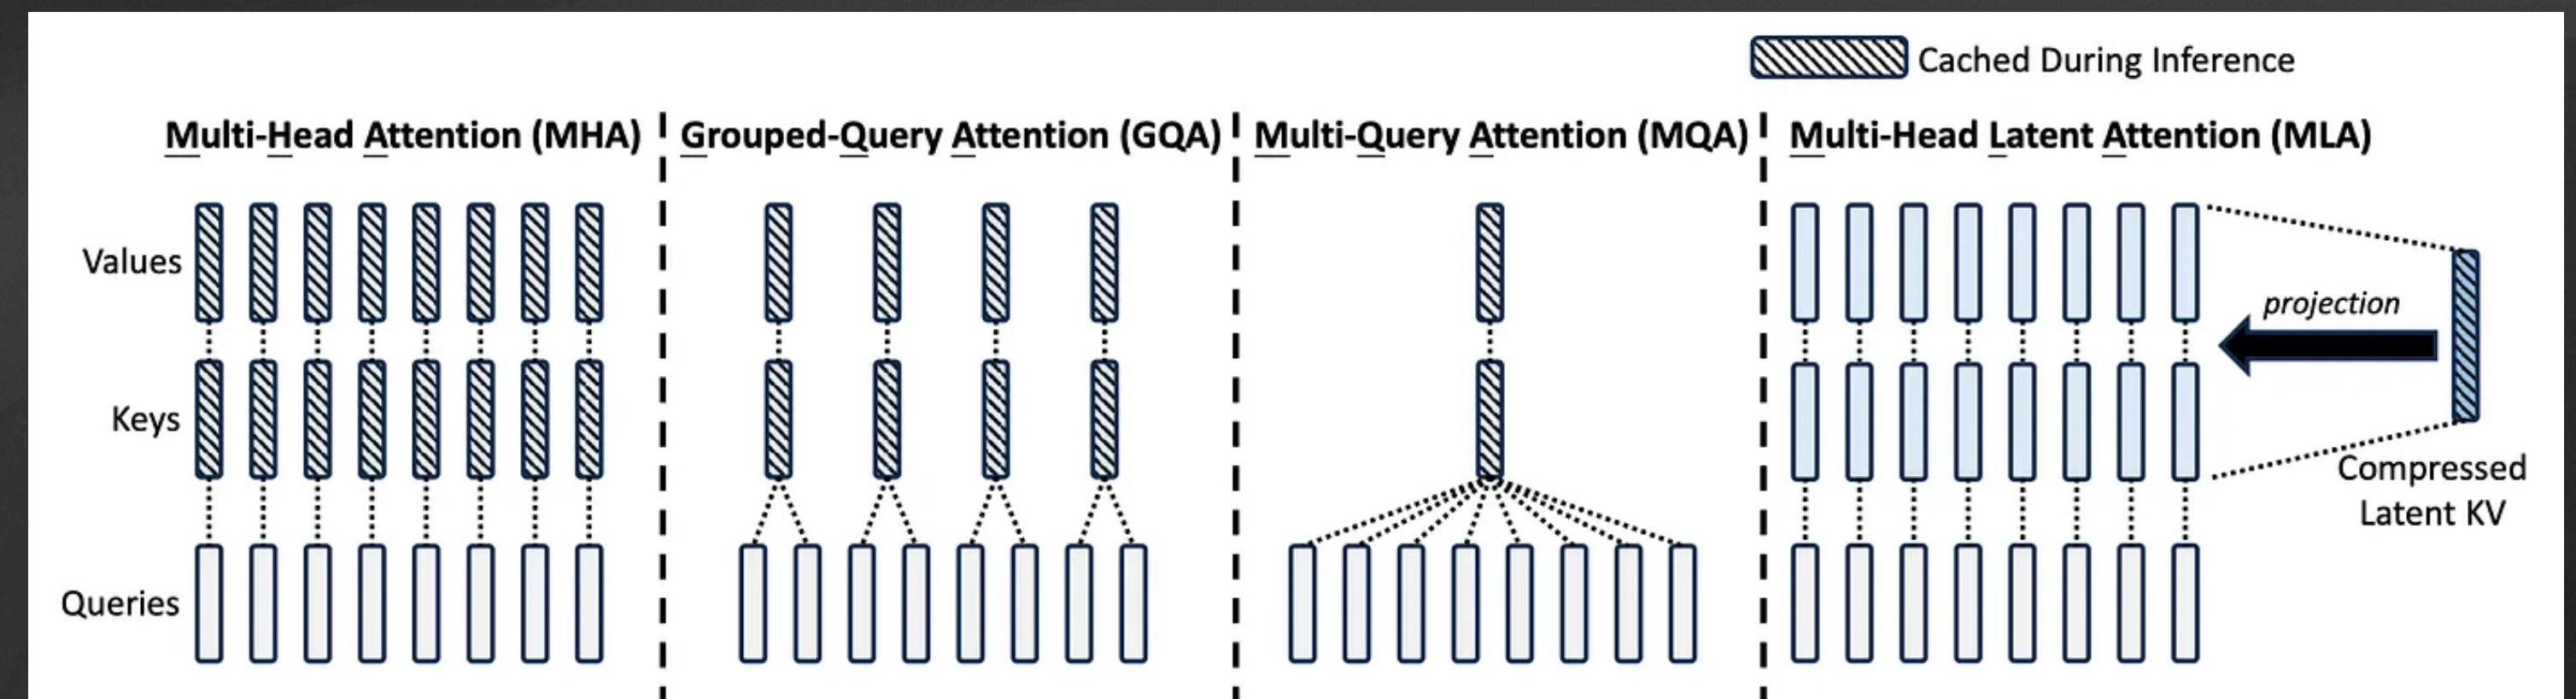


Figure 3 | Simplified illustration of Multi-Head Attention (MHA), Grouped-Query Attention (GQA), Multi-Query Attention (MQA), and Multi-head Latent Attention (MLA). Through jointly compressing the keys and values into a latent vector, MLA significantly reduces the KV cache during inference.

Model Improvement Techniques in Code

Assignment #1

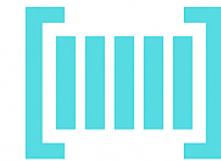
Your own dataset training

Dataset Search

<https://datasetsearch.research.google.com/>

<https://www.kaggle.com/datasets>

<https://paperswithcode.com/datasets>



Papers With Code

kaggle

Assignment #1

Your own dataset training

Submit a **Jupyter Notebook(.ipynb)** that demonstrates:

 **Step 1: Custom Model Declaration**

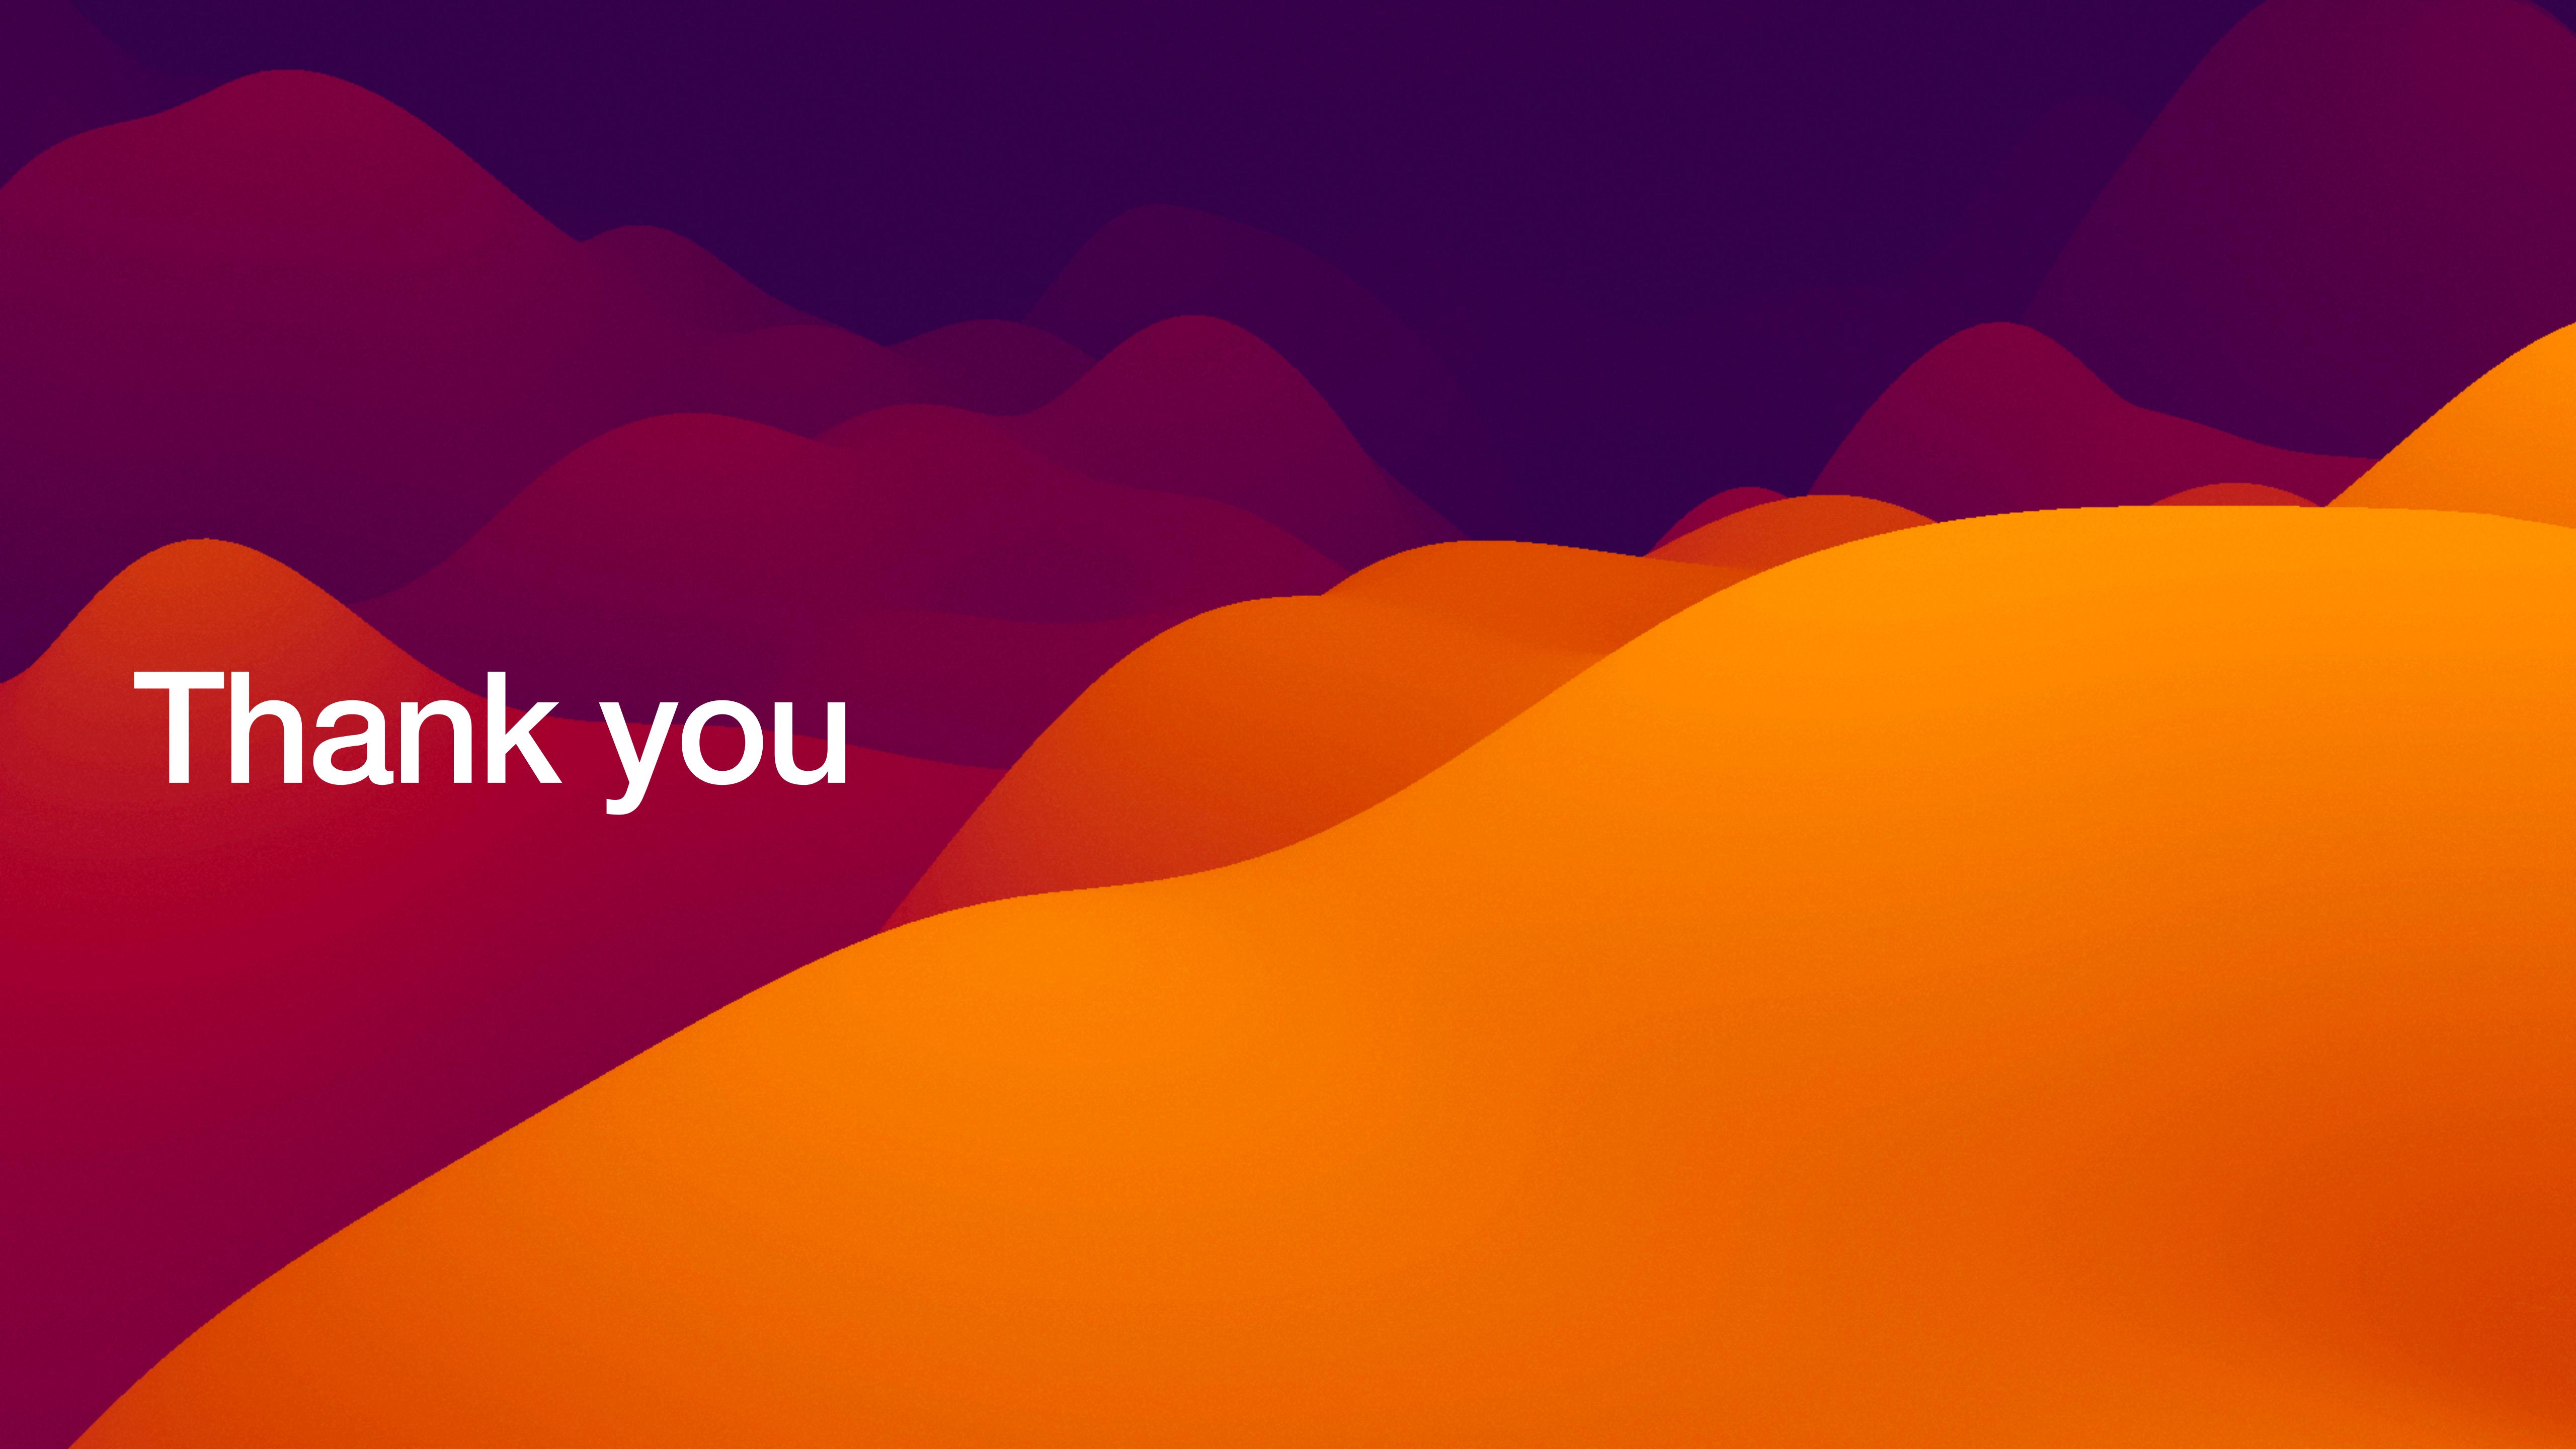
Define your custom model class with proper architecture initialization and forward pass implementation.

 **Step 2: Checkpoint Loading**

Load the saved model weights from your checkpoint file into the custom model instance.

 **Step 3: Dataset & Evaluation**

Load the specified dataset and perform model evaluation to demonstrate the trained model's performance.

The background features a minimalist design with abstract, undulating shapes. The lower half is filled with large, soft-edged waves in a bright orange color. Above this, the upper half consists of similar waves in a deep, saturated purple. The overall effect is one of calmness and modernity.

Thank you

Summary

Overall Code Structure

